

Computer Systems - CheatSheet

IN BA4 - Katerina Argyraki

Notes by Ali EL AZDI

A Computer Systems Cheatsheet has been authorized for the upcoming exam, and I'm sharing a copy of mine for anyone interested. It provides a concise summary of the key concepts and techniques covered in the course. For any updates or suggestions, feel free to reach out to me on Telegram at [elazdi_al](https://t.me/elazdi_al) or via EPFL email at ali.elazdi@epfl.ch.

April 3rd, 2025


Program. Passive entity. A sequence of instructions stored in a file, not currently executing.

- **Storage:** Stored on persistent storage (e.g., disk). Loadable into main memory by the OS.
- **Access:** Static file. Contains instructions, data, and metadata used during execution.

Process. Active execution of a program. Managed by the OS. Each process has its own isolated virtual address space.

- **Storage:** Occupies main memory during execution.
- **Memory Image:** Text (code), data, heap, and one stack per thread.
- **Process ID (PID),** assigned and tracked by the OS, unique system-wide.
- **Status:** OS tracks current status (e.g., running, waiting, terminated).
- **Virtualization:** Each process has the illusion of exclusive access to memory.
- **OS-allocated Resources:** File descriptors, sockets, I/O handles, etc.
- **Pointer to Page Table.** The **Kernel** tracks the process-specific page table pointer.
- **Access:** Operates in isolated virtual memory. No direct access to other processes. Interacts with hardware via system calls. OS handles scheduling and resource management.

If a process has a single thread, we may refer to the *thread's CPU context* as the *process's CPU context*.



Limited Direct Execution. OS design allowing user programs to execute instructions directly on the CPU, with restrictions.

- **Goal:** Maximize performance while maintaining control and protection.
- **User Code Execution:** CPU runs user programs natively (not emulated) in **user mode**.
- **OS Control:** OS retains control over hardware via **privileged instructions** and mode switching.

User Mode vs Kernel Mode. Two CPU execution modes stored as a state bit in a protected CPU register (0 = user mode, 1 = kernel mode) controlling access to hardware and instructions.

- **User Mode:** Restricted. User code cannot execute privileged instructions or directly access hardware/memory management.
- **Kernel Mode:** Mode in which is ran a central part of the operating system, the **Kernel**.
- **Switching:** A transition from user mode to kernel mode is triggered by:
 - System Calls (e.g., file I/O, memory allocation)
 - Hardware Interrupts (keyboard, timer,...)
 - Software Traps (divide-by-zero, invalid memory access),
If an exception happens in kernel mode, an internal routine is called but no mode switch occurs.

After handling the event, the OS switches the CPU back to user mode to resume application execution.

Privileged Instructions. CPU instructions that can only execute in **kernel mode**. If attempted in user mode, the CPU raises a trap to the OS. Prevent user programs from interfering with other processes or the OS.

To do something that requires high privilege, a thread makes a syscall, invoking the kernel.

Exception/trap - synchronous signal. The CPU itself raises an exception/trap.

Interrupt - asynchronous signal. Some external entity raises an interrupt. If an external entity needs attention, it raises an **interrupt**, invokes the kernel.

Even if nothing external needs attention, the timer interrupt regularly invokes the kernel to run the OS scheduler.

fork() — Clone the current process


1. The operating system creates a new process (child) by duplicating the calling process.
2. The child initially shares all memory pages with the parent:
 - Pages are marked read-only and shared (Copy-On-Write).
 - If either process attempts to write to a page, it is copied privately for that process.
3. The child also receives:
 - Duplicated file descriptors (pointing to the same open files).
 - Identical program counter and stack pointer.
4. Returns:
 - In the parent process: the PID of the child.
 - In the child process: 0.
5. Both processes resume execution at the instruction following **fork()**.

exec() — Mutate the child into a new program

- Typically called by the child process immediately after **fork()**.
- Replaces the child's memory space with a new program image.
- Stack, heap, code, and data segments are replaced.
- If successful: the new program starts execution from **main()**.
- If failed: **exec()** returns -1, child continues old code.

Thread. The smallest unit of CPU execution within a process. Multiple threads can exist within one process.

- **Memory:** Shares the parent process's virtual address space:
 - Shared: Code (text), heap, global/static data, open files.
 - Private: Each thread has its own stack (for local variables and function calls).
- **Execution context (per thread), values of all the CPU registers at the last moment the thread was running:**
 - **Instruction Pointer (IP) / Program Counter (PC):** Physical register pointing to the next instruction. Private per thread.
 - **Stack Pointer (SP):** Points to the top of the thread's private stack.
 - **General-purpose Registers:** Include temporary registers (e.g., RAX, RBX), status, and flags. **All private per thread.**
 - **Thread ID (TID).** The Thread ID is unique within a process, but not necessarily system-wide, assigned and tracked by the OS.
 - **Status:** OS tracks each thread's status (Running, Ready, Blocked).
 - **Virtualization:** Each thread has the illusion it exclusively occupies the CPU.
 - **Context switching:** OS saves/restores full register set (IP, SP, general-purpose regs). The CPU switches from one thread to another. Thread switches are faster than process switches. It saves to memory the CPU context of the current thread; it restores from memory the CPU context of the new thread.
 - **Thread management:** A main thread may create and manage others. In some models, a dedicated manager thread exists.
 - **Managed by OS:** The OS kernel schedules threads individually and tracks them.



OS Components.

- **Kernel (core of the OS):** Loaded at boot, runs in **kernel mode**, and manages hardware, memory, processes, and system calls. It is the only part of the OS that runs in kernel mode.
Three important routines inside of the kernel:
 - **Interrupt Handler or Interrupt Service Routine (ISR):** Handles interrupts from hardware devices.
 - **Exception / Trap Handler:** Handles exceptions and traps from user programs.
 - **System Call Handler:** Handles system calls from user programs.
- **The three routines use two tables (Trap/Interrupt Table and System Call Table) to resolve the event type to a specific memory address of the handler routine managing the event.**
- **Loader (part of the OS):** Prepares executables to run:
 - Loads program code/data into memory.
 - Maps required libraries (e.g., libc).
 - Sets up the process stack and heap.
 - Places command-line arguments and environment variables on the stack.
 - Sets the **%rsp** (stack pointer) and **%rdi** (argc), **%rsi** (argv) registers so the program can access arguments.
 - *Because the kernel manages processes and initiates execution, it is responsible for placing arguments in registers so the loader (running in user mode) can access them.*
 - Jumps to the program's entry point to begin execution in **user mode**.
- **User-level Programs:** Applications like shells, editors, browsers, etc. These are **part of the OS** but run entirely in **user mode**, relying on system calls to request kernel services.
- **System Libraries:** Shared libraries (e.g., libc) used by user programs. Loaded and linked by the loader, but executed in user mode. Provide wrappers around system calls.

wait() — Wait for a child process to terminate


- Blocks the calling process until one of its child processes terminates.
- Returns the PID of the terminated child.
- If no child processes exist, it waits indefinitely.

Memory Management Unit (MMU) Address Translation.

Hardware component managed by the operating system that translates hardware addresses into physical addresses using per-process page tables. Each process has its own page table, and the currently active one is referenced in a special CPU register.

Translation Lookaside Buffer (TLB).

A small, fast **cache** inside the MMU that stores recent **virtual-to-physical page mappings**. The TLB reduces the need to access the page table on every memory reference, significantly speeding up address translation.



Virtual Address Translation:

1. The CPU executes an instruction (e.g., load, store, or fetch) that references a memory location using a **virtual address**.
2. The MMU extracts the **virtual page number (VPN)** and **page offset** from the virtual address.
3. The MMU first checks the **TLB** for a cached translation of the VPN.
4. **If the translation is found in the TLB (TLB hit):**
 - (a) The corresponding **physical frame number** is retrieved directly.
5. **If the translation is not found (TLB miss → page table walk):**
 - (a) The MMU uses the active **page table** to look up the VPN and obtain the PFN.
 - (b) The new mapping is inserted into the TLB for future accesses.
6. **If the page table entry is invalid or present = 0 (page fault):**
 - (a) The CPU triggers a **trap** into the OS kernel.
 - (b) The **page fault handler** in the OS's memory management subsystem is invoked.
 - (c) The handler checks whether the page is:
 - i. **Never allocated before** → allocate a new physical page and zero-initialize it.
 - ii. **Swapped out to disk** →
 - A. Check the **swap cache** for the page.
 - B. If found in the swap cache, use the cached page directly.
 - C. If not in the cache, read the page from the **swap space** on disk into a free physical frame, and insert it into the swap cache.
 - (d) The newly loaded or allocated page is mapped in the page table, and **present** is set to 1.
 - (e) The TLB entry for the VPN is updated if needed.
 - (f) The faulting instruction is retried.
7. The final **physical address** is then used to **access RAM**.

Swapping

Transfer of memory pages between physical RAM and disk-based swap space to free up RAM.

- **Swap Space** Disk region reserved for evicted pages; holds non-resident memory to extend usable RAM.

- **Swap Cache** In-memory buffer of recently swapped-out pages; enables fast lookup and avoids redundant disk I/O during swap-in.

-Swap-Outs:

1. The OS detects that the number of free physical pages has fallen below a predefined threshold.
2. A background kernel thread is triggered to reclaim memory.
3. The kernel scans memory to identify candidate pages for eviction, typically using an aging or Least Recently Used algorithm.
4. A candidate page is selected if:
 - (a) It is not currently in use (i.e., not recently accessed),
 - (b) It is not locked, pinned, or shared with kernel-critical structures.
5. If the selected page is dirty (i.e., has been modified), its contents are written to the swap space on disk.
6. The page is inserted into the **swap cache** so it can be quickly retrieved if needed again.
7. The page table entry is updated: **present** = 0, and the swap location is recorded.
8. The physical frame is freed and returned to the pool of available memory.

CPU Cache.

The **CPU cache** is a small, fast memory located on or near the processor that stores copies of frequently accessed data from main memory (RAM). Its purpose is to reduce memory access latency and improve overall performance by exploiting **temporal** and **spatial locality**.

Cache Levels:

- **L1 Cache (Level 1):** The smallest and fastest cache, located right on the CPU core. 64 KB, but takes $< 1nsec$ to access.
- **L2 Cache (Level 2):** Larger and slightly slower than L1, still located on the CPU core. 256-512 KB, takes $< 4nsec$ to access.
- **L3 Cache (Level 3):** Shared among multiple cores, larger but slower than L1 and L2. Located on the CPU die. 6-32 MB, takes 10s of $nsec$ to access.

Hierarchy Behavior:

Caches form a **hierarchy**: the CPU checks L1 first, then L2, then L3, and finally main memory if needed. This optimizes for latency and hit rate.

Paging.


Paging divides virtual memory into fixed-size **pages**, mapped to physical frames via a **page table**. It avoids **ext. fragmentation**, but may cause **internal fragmentation** if a page is only partially used.

Address Representation.

Virtual Address Size is not always equal to Physical Address Size!

Let m be the number of bits in the **virtual address** (i.e., $m = \log_2(\text{virtual address space size})$).

This also corresponds to the number of bits required to uniquely address every byte in the virtual memory. If the virtual memory consists of 2^k pages and each page is 2^n bytes, then the total virtual address space is $2^k \times 2^n = 2^m$ bytes, and thus $m = k + n$.



A virtual address is split into:

- **Offset** (n bits): identifies a byte within a page of size 2^n bytes, where $n = \log_2(\text{page size})$.
- **Virtual Page Number (VPN)** ($m - n$ bits): selects the page entry from the page table.

Page Table.

A process's **page table** maps each **Virtual Page Number (VPN)** to a **Physical Frame Number (PFN)**. The page table is indexed by the VPN and stores **Page Table Entries (PTEs)**.

Linear Page Table.

A single-level table with 2^{m-n} entries. Each VPN directly indexes a PTE. Simple but potentially large: $\Rightarrow$ Table size = $2^{m-n} \times \text{PTE size}$ (may span multiple memory pages).

PTE Content (Typical Metadata Bits):

- **Present bit (P):** Valid address translation exists
- **Protection bits (R/W/X):** Read, write, execute permissions
- **User/Supervisor (U/S):** Access control (user vs. kernel)
- **Dirty bit (D):** Set if the page has been modified
- **Access/Reference bit (A):** Set on access; used in replacement policies

Multi-Level Page Tables.

Used to reduce memory overhead. Split the VPN into multiple parts to form a tree-like hierarchy. The **offset** remains the last $n = \log_2(\text{page size})$ bits of the address.

General Breakdown (Multi-Level Paging):

$$\text{Virtual Address Size} = m = \sum k_i + n$$

where:

- $n = \log_2(\text{page size})$: offset bits
- k_i : number of bits used at each level of the page table

Two-Level Page Table:

$$\text{VPN} = \underbrace{\text{Level 1 index}}_{k_1} \parallel \underbrace{\text{Level 2 index}}_{k_2}, \quad k_1 + k_2 = m - n$$

- Level 1 index (k_1 bits): selects the **page directory entry (PDE)**
- Level 2 index (k_2 bits): selects the **page table entry (PTE)** from the second-level table

Three-Level Page Table:

$$\text{VPN} = \underbrace{\text{L1 index}}_{k_1} \parallel \underbrace{\text{L2 index}}_{k_2} \parallel \underbrace{\text{L3 index}}_{k_3}, \quad k_1 + k_2 + k_3 = m - n$$

- Level 1 index (k_1 bits): selects **first-level page directory**
- Level 2 index (k_2 bits): selects **second-level directory/table**
- Level 3 index (k_3 bits): selects the final **PTE**

Bit Allocation in Multi-Level Page Tables:

The VPN portion ($m - n$ bits) is divided into L parts ($k_1 + k_2 + \dots + k_L = m - n$), one per level.

- **Bit division is not necessarily even.** Systems may assign more bits to higher levels.
- If uneven, **last levels (closer to the leaf)** typically receive **fewer bits** (fewer entries).
- Each level i contains 2^{k_i} entries, indexing the next level or the final PTE.

Segmentation.

Segmentation divides memory into variable-sized **logical segments** (e.g., code, data, stack), each with a base and limit. It aligns with program structure and supports segment-level protection, but suffers from **ext. fragmentation** due to variable-sized allocations.

Block Cache. In-memory cache of disk blocks

- **Good for reads** — avoids repeated disk access
- **Limited for writes** — consistency requires flushing
- **Write-through** — data written to both cache and disk immediately (slower, but consistent)
- **Write-back** — data stays in cache, flushed later (faster, but risk of data loss on crash)
- **fsync()** — forces flush of dirty blocks to disk

Consistency Update Problem.
File system metadata may become inconsistent due to crashes during updates.

File System Checker (fsck).
Tool to scan and repair on-disk metadata inconsistencies.

Cons: Functionality - fix not always obvious or correct,
Performance - slow; may take hours.

fsck Fix Examples.

- **Link Count Inconsistency** — problem: inode's link count $\neq$ number of directory entries pointing to it; fix: correct link count to match actual references.
- **Lost Inodes** — problem: inode has link count > 0 but no directory entry points to it; fix: move file to **lost+found** for recovery.
- **Data Bitmap Errors** — problem: inode uses a block, but bitmap marks it free; fix: set corresponding bitmap bit to "used".
- **Duplicate Pointers** — problem: two inodes point to same data block; fix: duplicate the block and update one inode to preserve both files.
- **Invalid Pointers** — problem: inode points to a block beyond partition size; fix: remove invalid pointer to prevent access.

Journaling.
Crash-consistency technique. Metadata updates are first written to a log (journal) before applying to the main file system

Pros. Fast crash recovery (no full scan); Metadata consistency guaranteed; Safer than fsck in most cases.

How It Works.

1. Group changes into a **transaction**
2. Write to journal: TxBegin | changes | TxEnd | Valid
3. Valid block = transaction committed
4. Apply changes to file system (check-point)
5. Clear transaction from journal

I/O Interrupt.
Device triggers an interrupt when it needs service.

Pro: efficient for unpredictable events.

Con: high overhead per interrupt.

LiveLock.
System busy handling I/O (e.g., interrupts) but makes no progress in real work.

Real-World Strategy.

- Use **interrupts** for overlap (slow devices)
- Use **polling** for short bursts, small data, high performance
- **Coalescing:** delay, batch multiple responses for efficiency

DMA (Direct Memory Access).
CPU tells the device **where data is.**

DMA Transfer Workflow. (device to memory transfer.)

1. The device driver receives an instruction to transfer disk data to a buffer at address X.
2. The driver commands the disk controller to transfer C bytes from disk to the buffer at address X.
3. The disk controller initiates the DMA transfer operation.
4. The disk controller sends each byte to the DMA controller.
5. The DMA controller transfers bytes to buffer X, incrementing the memory address and decrementing C until C = 0.
6. When C = 0, the DMA controller interrupts the CPU to signal completion of the transfer.

ls -l command output.

pe	Usr	Grp	Oth	Ln	Owner	Group	Size	Mon	Day	Time/Year	Name
/-/1/	ppp	ppp	ppp	[n]	[username]	[groupname]	[bytes]	[Mon]	[DD]	[HH:MM / YYYY]	[filename]

Permissions Breakdown:

Type — File type:
- - = regular file
- d = directory
- l = symbolic link
Usr, Grp, Oth — Permission triplets for:
- r = read
- w = write
- x = execute
- - = permission not granted

- Ln — Number of hard links
- Owner — Username of the file owner
- Group — Group associated with the file
- Size — File size in bytes
- Mon, Day, Time/Year — Last modification timestamp
- Name — File or directory name

malloc(size_t size) | Allocates size bytes. Returns void* to uninitialized memory.

Return: pointer to memory or NULL on failure

calloc(size_t nmemb, size_t size) | Allocates space for nmemb elements of size bytes each. Memory is zero-initialized.

Return: pointer to zeroed memory or NULL

memset(void *ptr, int value, size_t num) | Sets the first num bytes of ptr to (unsigned char)value.

Use: initialize or reset memory

realloc(void *ptr, size_t new_size) | Resizes previously allocated block to new_size bytes.

Return: pointer to new block (contents preserved), or NULL


Disk Storage
Goal: fast, reliable, affordable persistent data access. Must return what was written-quickly and without loss.

RAID (Redundant Array of Inexpensive Disks)
Combines multiple physical disks into one logical unit for performance and fault tolerance. Allows for a higher throughput and reliability.

RAID 0 (Striping)
data split across $N \geq 2$ disks
+ High throughput, full capacity
- No redundancy
Read: from corresponding disk
1 disk fails: all data lost

RAID 5 (Striping + Parity) — block-level striping + distributed parity across $N \geq 3$ disks
+ Efficient redundancy, survives 1-disk failure
- Writes = read + modify parity (slow)
Read: from data disk if intact
1 disk fails: reconstruct via XOR:
Missing Block = Parity $\oplus$ $\bigoplus_{i \neq \text{failed}} \text{Data}_i$
Or define parity as:
Parity = $\bigoplus_{i=1}^{N-1} \text{Data}_i$

RAID 1 (Mirroring)
data duplicated on 2 disks
+ Survives 1-disk failure, fast reads
- $2\times$ storage cost
Read: from either copy (load-balanced)
1 disk fails: use remaining mirror



Context Switch
Triggered by *timer interrupt* or blocking events. Used to enforce fairness and preempt misbehaving threads.

Procedure.

1. **Save state** of current thread (PC, registers, etc.) to PCB
2. **Choose next thread** via scheduler
3. **Restore state** of selected thread from its PCB
4. **Resume execution** via **return-from-trap**

Timer interrupts ensure control returns to the OS periodically, allowing preemption and preventing CPU hogging.

Scheduling Policy — strategy for choosing the next thread

Scheduling Metrics

- **CPU Utilization:** % time CPU is busy
- **Turnaround Time:** completion time – submission time
- **Response Time:** time from submission to first response (eg. submission to first time a thread is scheduled)

FIFO Scheduling (First-In, First-Out)
Non-preemptive (waits for completion before scheduling next thread). Runs threads in order of arrival.

1. **Enqueue** new threads at the tail of the ready queue
2. **Select** the thread at the head of the queue
3. **Run** it to completion or blocking
4. **Repeat** with the next thread in queue

Pros: Simple, fair (by arrival order), minimal overhead
Cons: Poor responsiveness, suffers from *convoy effect* (short jobs wait behind long ones)

STCF Scheduling (Shortest Time to Completion First)
Preemptive SJF. Always runs the thread with the least remaining time.

1. **Track** remaining time for all ready/running threads
2. **On arrival** of a new thread, compare its time to current thread
3. **Preempt** if new thread has shorter remaining time
4. **Run** the shortest job until completion or preemption

Pros Optimal average turnaround time under preemption
Cons Requires accurate time estimates, may cause starvation of longer jobs, bad Response Time

SJF Scheduling (Shortest Job First)
Non-preemptive. Picks the thread with the shortest remaining execution time.

1. **Estimate** or know job lengths in advance
2. **Select** the shortest job from the ready queue
3. **Run** it to completion or blocking
4. **Repeat** with the next shortest job

Pros Minimizes average turnaround time (optimal under perfect knowledge)
Cons Requires job length estimates, risk of starvation for long jobs

MLFQ Scheduling (Multi-Level Feedback Queue)
Preemptive. Dynamically adjusts thread priority based on behavior.

Rules

1. If priority(A) > priority(B), **A runs**
2. If priority(A) = priority(B), **A, B run in Round Robin**
3. **New threads start at top priority**
4. If a thread uses up its time slice, **demote its priority**
5. Periodically **boost all threads to top priority** (prevents starvation)

Procedure

1. Maintain multiple ready queues by priority level
2. Insert new or boosted threads at highest priority
3. Select thread from highest non-empty queue
4. Run using Round Robin within queue
5. **Demote if thread uses full time slice; boost periodically !!!!**


Round Robin Scheduling
Preemptive. Each thread gets a fixed time slice, cycling through the ready queue.

1. **Enqueue** threads in arrival order
2. **Run** the thread at the head for one quantum
3. **Preempt** if not finished; move to tail of queue
4. **Repeat** with the next thread in queue

Pros Fair, responsive, avoids starvation
Cons High context switch overhead if time slice is too small; poor for short jobs if time slice is too large

Generic Pointers.

void apply(void* array, size_t len, size_t elem_size, void (*func)(void*)) {
for (size_t i = 0; i < len; i++) {
void* elem = (char*)array + i * elem_size;
func(elem);
}
}



Recursive Struct.

```
typedef struct Node {  
    int value;  
    struct Node* next;  
} Node;
```

Network Addressing & Web Protocols:

- IP Address:** 32-bit (IPv4) or 128-bit (IPv6) unique identifier for network devices. Format: dotted decimal (e.g., 192.168.1.1) or hexadecimal notation.
- Port Number:** 16-bit identifier (0-65535) for specific application/service on a host. Well-known ports: HTTP=80, HTTPS=443, FTP=21, SSH=22.

HTTP (HyperText Transfer Protocol): Stateless application-layer protocol for web communication.

- Stateless Protocol:** Server doesn't maintain client state between requests. Each request is independent and self-contained.
- Web Objects:** Files identified by URLs (HTML pages, images, videos, CSS, JavaScript, etc.). Each object has a unique URL.
- Web Page:** Base HTML file + referenced objects (images, stylesheets, scripts). Single page may require multiple HTTP requests.

HTTP Methods:

GET: Retrieve resource from server.

Safe, idempotent.

HEAD: Like GET but returns only headers (no body). Used for metadata.

POST: Submit data to server (forms, uploads). Non-idempotent.

PUT: Update/create resource. Idempotent.

DELETE: Remove resource. Idempotent.

Cookies: Small data stored by browser, sent with requests to maintain state across stateless HTTP. Used for authentication, preferences, tracking.

- Web Caching:** Store frequently accessed objects closer to client to reduce latency and server load. Types: browser cache, proxy cache, and **Edge Caches** (part of a CDN) which are geographically distributed.

- Max-Age:** Cache-Control directive specifying how long (in seconds) a resource can be cached before becoming stale. Example: **max-age=3600** (1 hour).

DNS (Domain Name System): Hierarchical distributed naming system that translates human-readable domain names to IP addresses.

- DNS Name:** Hierarchical structure (e.g., www.epfl.ch). Read right-to-left: .ch (TLD), epfl (2nd level), www (subdomain).

- Reserved Port:** UDP/TCP port 53. UDP for standard queries, TCP for zone transfers and large responses.

DNS Server Hierarchy:

- Root DNS Servers:** 13 logical servers (A-M) managing top-level domains
- TLD DNS Servers:** Manage specific top-level domains (.com, .org, .ch, etc.)
- Authoritative DNS Servers:** Store actual DNS records for specific domains
- TTL (Time To Live):** Specifies how long DNS records can be cached (in seconds). Balances performance vs freshness. Common values: 300s (5min) to 86400s (24h).

- Protocol:** Primarily UDP for speed. TCP used for zone transfers or when response exceeds 512 bytes.


Client-side Syscalls (UDP):

- socket():** Creates a new UDP socket endpoint.
- bind():** (Optional) Assigns a local port to the socket. If not called, an ephemeral port is assigned by the kernel on first 'sendto()'.
- sendto():** Sends the DNS query to the server's IP and port 53.
- recvfrom():** Waits to receive the DNS response from the server.

DNS Resolution Approaches:


Iterative Query:

- Client queries local DNS
- Local DNS queries root, gets TLD referral
- Local DNS queries TLD, gets authoritative referral
- Local DNS queries authoritative, gets answer
- Each server returns **referral** or answer



Recursive Query:

- Client queries local DNS
- Local DNS recursively queries on behalf of client
- Each server in chain does the work
- Final answer returned through chain
- Client only sees final result



DNS client

DNS Structure:

- ISP (Internet Service Provider):** Company that provides Internet access.
 - Tier-1 ISP:** Global network, peers with all other Tier-1s. The backbone.
 - Regional ISP:** Connects to one or more Tier-1 ISPs.
 - Access ISP (Last-mile):** Connects end-users (homes, businesses).
- IXP (Internet Exchange Point):** Physical infrastructure where multiple ISPs connect to exchange traffic directly, avoiding Tier-1 transit.

Packet Switching & Network Delays:

- Packet Switch:** Network device that forwards packets using store-and-forward switching. Receives entire packet before forwarding.
- Transmission Delay:** Time to push packet onto link. $d_{trans} = \frac{L}{R}$, where L is the packet size (bits), R is the transmission rate (bps).
- Propagation Delay:** Time for signal to travel link. $d_{prop} = \frac{d}{s}$, where d is the distance, s is the signal speed.
- Queuing Delay:** Time packet waits in output queue before transmission. Depends on congestion level.
- Total Packet Delay:** $d_{total} = d_{proc} + d_{queue} + d_{trans} + d_{prop}$, where d_{proc} is the processing delay.
- Packet Loss:** Occurs when packet arrives at full queue buffer. Packet is dropped.
- Traffic Analysis:**
 - Arrival Rate (λ):** Average rate at which packets arrive.
 - Departure Rate (μ):** Rate at which packets are served/transmitted.
 - If $\lambda > \mu$: Queue grows indefinitely, delay increases, eventually packet loss.
 - If $\lambda < \mu$: Stable system, finite queuing delay.
 - If $\lambda = \mu$: System at capacity, high sensitivity to bursts.
- Average Throughput:** $\min(\text{transmission rates along path})$ - the bottleneck link determines end-to-end throughput.

UDP (User Datagram Protocol): Connectionless, unreliable transport protocol.

- Characteristics:** No connection establishment, no reliability guarantees, no flow/congestion control, minimal overhead.

Header Structure (8 bytes total):

- Source Port: 16 bits
- Destination Port: 16 bits
- Length: 16 bits (UDP header + data length)
- Checksum: 16 bits (optional in IPv4, mandatory in IPv6)

Data Encapsulation (5-Layer Model):

- Application Layer:** Creates message. (e.g., HTTP GET)
- Transport Layer:** Adds TCP/UDP header (ports) → **Segment**.
- Network Layer:** Adds IP header (source/dest IP) → **Datagram**.
- Link Layer:** Adds MAC header (e.g., Ethernet) → **Frame**. Managed by the **Network Interface (NIC)**, which is the hardware connecting a device to a network.
- Physical Layer:** Transmits frame as bits on the medium.

TCP Protocol Essentials. Transmission Control Protocol provides reliable, ordered delivery over unreliable networks.

- Three-Way Handshake:** Connection establishment process.

- Step 1:** Client sends SYN with initial sequence number (ISN)
- Step 2:** Server responds with SYN-ACK (server ISN + ACK of client ISN+1)
- Step 3:** Client sends ACK (acknowledges server ISN+1)

- Sequence Number (SEQ):** 32-bit field identifying byte position in data stream. Initial value chosen randomly. Each byte of data increments SEQ by 1. Used for ordering and duplicate detection.

- Acknowledgment Number (ACK):** 32-bit field indicating next expected sequence number from sender.

- Congestion Window (cwnd):** TCP sender's estimate of how many bytes can be sent without causing network congestion, dynamically adjusted based on network conditions.

- Purpose:** Controls sending rate to prevent network overload and packet loss due to buffer overflow at **intermediate routers**.

- Relationship:** $\text{Effective Window} = \min(\text{cwnd}, \text{rwnd})$ where rwnd = receiver advertised window.

$$\text{ACK} = \text{last_received_SEQ} + 1$$

Cumulative acknowledgment (acknowledges all data up to ACK-1).

- Maximum Segment Size (MSS):** Largest amount of data bytes that TCP can send in a single segment, excluding TCP and IP headers.

$$\text{MSS} = \text{MTU} - \text{IP_header_size} - \text{TCP_header_size}$$

where MTU = Maximum Transmission Unit of the network path (typically 1500 bytes for Ethernet), IP header = 20 bytes (without options), TCP header = 20 bytes (without options). Standard Ethernet MSS = 1460 bytes. MSS is negotiated during TCP connection establishment via MSS option in SYN segments.

- Timeout Calculation:** TCP adaptively calculates retransmission timeout to balance quick recovery from lost packets with avoiding unnecessary retransmissions due to network delay variations.

- SampleRTT:** Measured time for a specific segment to be sent and acknowledged (not retransmitted segments)

- EstimatedRTT:** Exponentially weighted moving average of RTT samples to smooth out fluctuations

- DevRTT:** Estimate of RTT variation to account for network jitter and delay variability

$$\text{EstimatedRTT} = (1 - \alpha) \times \text{EstimatedRTT} + \alpha \times \text{SampleRTT}$$

where α = smoothing factor for RTT estimation (typically 0.125)

$$\text{DevRTT} = (1 - \beta) \times \text{DevRTT} + \beta \times |\text{SampleRTT} - \text{EstimatedRTT}|$$

where β = smoothing factor for deviation estimation (typically 0.25)

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 \times \text{DevRTT}$$

- Checksum:** 16-bit Internet checksum. Computed over TCP header, data, and pseudo-header (source IP, dest IP, protocol, TCP length). Uses 1's complement arithmetic.

- TCP Congestion Control (Tahoe & Reno):** Algorithms that dynamically adjust sending rate to prevent network congestion while maximizing throughput.

Common Phases (Both Algorithms):

- Slow Start:** $\text{cwnd} = \text{cwnd} + \text{MSS}$ for each ACK. Exponential growth until $\text{cwnd} \geq \text{ssthresh}$.

- Congestion Avoidance:** $\text{cwnd} = \text{cwnd} + \frac{\text{MSS}^2}{\text{cwnd}}$ for each ACK. Linear growth.

Fast Retransmit vs No Fast Retransmit:

- Enabled:** Upon 3 duplicate ACKs, immediately retransmit lost segment. Faster recovery, reduced latency.
- Disabled:** Must wait for timeout expiration to detect and retransmit lost segments. Slower recovery, higher latency.

Loss Recovery (Key Difference):

Reno (Adaptive):

- Timeout:** $\text{ssthresh} = \text{cwnd}/2$, $\text{cwnd} = \text{MSS}$ (slow start)
- 3 Dup ACKs:** $\text{ssthresh} = \text{cwnd}/2$, $\text{cwnd} = \text{ssthresh} + 3 \times \text{MSS}$ (fast recovery)
- Behavior:** Distinguishes mild vs severe congestion

Tahoe (Conservative):

- Timeout / 3 Dup ACKs:** $\text{ssthresh} = \text{cwnd}/2$, $\text{cwnd} = \text{MSS}$ (slow start)
- Behavior:** Always returns to slow start after ANY loss

- Summary:** Tahoe always returns to slow start after any loss (conservative). Reno uses fast recovery for mild congestion (3 dup ACKs) but slow start for severe congestion (timeout).